

Introduction to Machine Learning

Homework 3

Due: Tuesday, 10/29, 5pm

Instructor:

Dr. Tom Arodz




The problem

- Implement and test:
 - Linear classifier for 10 classes
 - With cross entropy loss (over softmax)
 - With L1 regularization (see below)
 - Trained using gradient descent
- Use PyTorch library


Dataset

- MNIST: a 10-class classification problem
 - 28x28 gray-scale images of handwritten digits

Label: 7



Label: 5



- Training set of 50,000 samples (5,000 per class)
- Test set of 10,000 samples (1,000 per class)

```
import torch
from torchvision import datasets, transforms
import matplotlib.pyplot as plt
```

```
def show_mnist_image(index):
    # Load MNIST dataset
    mnist_data = datasets.MNIST(
        root='./data', train=True,
        download=True,
        transform=transforms.ToTensor())
```

```
# Get the image and label
image, label = mnist_data[index]
```

```
# Convert image from PyTorch tensor
# to numpy array and reshape
image = image.squeeze().numpy()
```

```
# Display the image
plt.imshow(image, cmap='gray')
plt.title(f"Label: {label}")
plt.axis('off')
plt.show()
```

```
show_mnist_image(42)
```



Your task

- Write a function to train a model on the dataset, and a function to test it on another dataset

```
def my_train(train_dataset)
    #your code
    return W, b
```

```
def my_test(W,b,test_dataset)
    #your code (error rate on 0-1 scale)
    return test_error_rate
```

- Your goal for the `my_train` function is to train the model to achieve high accuracy (at least 87.5% accuracy, i.e., 0.125 (or, 12.5%) error rate), while keeping some weights for redundant features at 0
 - You will need to select appropriate hyperparameter (e.g., learning rate, number of epochs, regularization strength)
- Your goal for the `my_test` function is to correctly calculate the error rate, using pytorch functions (incl. DataLoader)



L1 regularization

- L1 regularization is an approach similar to L2 regularization (weight decay)
- It aims at reducing the number of irrelevant features, by bringing their weights to 0
- $\text{Loss_with_regularization}(W) = \text{CrossEntropy}(W) + \lambda ||W||_1$
 - λ controls the “strength” of regularization, with $\lambda=0$ meaning no regularization, and high λ means strong regularization (more features will become 0)
- $||W||_1$ is L_1 -norm, defined as sum of absolute values of all the matrix cells
 - $||W||_1 = \sum_{ij} |W[i,j]|$



Your task – variable details

- ```
def my_train(train_dataset)
 #your code
 return W, b

def my_test(W,b,test_dataset)
 #your code (error rate on 0-1 scale)
 return test_error_rate
```
- W, b are be numpy arrays,
  - shape of W: (num\_classes, num\_features)
  - shape of b: (num\_classes,)
- test\_error\_rate is single number
- train\_dataset, test\_dataset are of type:  
**torchvision.datasets.mnist.MNIST**  
loaded as in the next slide



# Classification model and training

- Dataset loading that will be used in grading:

```
from torchvision import datasets, transforms
```

```
transform = transforms.Compose([
 transforms.ToTensor(),
 transforms.Normalize((0.1307,), (0.3081,))
])
```

```
train_dataset = datasets.MNIST(root='./data',
 train=True, download=True, transform=transform)
```

```
test_dataset = datasets.MNIST(root='./data',
 train=False, download=True, transform=transform)
```



# Model and objective function

---

- The classification model should be:
  - A linear model ( $W^T x + b$ )
    - shape of  $W$ : (num\_classes, num\_features)
    - shape of  $b$ : (num\_classes,)
    - num\_features should be 784  
(28x28 pixels, each one grayscale value)
  - It is ok to use nn.Linear, but observe the shape of their internal parameter (.weight and .bias) and convert the shape if it doesn't match the above specs for input/output of your functions my\_train and my\_test
- The objective function should be cross-entropy loss + L1 regularization
  - note whether it should take raw predictions (a.k.a. logits, from  $[-\infty, +\infty]$  range), or probabilities ( $[0, 1]$  range)





# Returning the Assignment

---

- Solution code should be written by you and you only (no web/book/friend/etc. code)
- Upload through Canvas/Gradescope
  - Similar to Homework 1 & 2
  - A single file with your two functions
    - Do not forget to do all the necessary imports
    - If your code doesn't "compile" or throws an exception, gradescope will fail, with 0 points
    - It is advisable to either delete any of your testing code, or "guard" it with:  
`if __name__ == "__main__":`