

Introduction to Machine Learning

Lecture 2

Instructor:
Dr. Tom Arodz



Machine Learning

- Machine learning:
 - Study of algorithms that
 - improve their performance P
 - at some task T
 - with experience E
 - We have a well-defined learning task: $\langle P, T, E \rangle$



Machine Learning

- Machine learning: study of algorithms that
 - improve their performance P
 - HOW TO MEASURE THE PERFORMANCE?
 - at some task T
 - WHAT IS THE TASK?
 - with experience E
 - HOW IS EXPERIENCE GAINED?






Supervised Machine Learning

- Algorithms for solving the learning problem:
<Performance, Task, Experience>
- **Supervised** Machine Learning/Classification:
 - Task:
 - Learn the ability to categorize objects:
to predict the class of object
based on some attributes/features of the object
 - Performance:
 - Measured as the accuracy (correctness) of predictions
for previously unseen objects
 - Experience:
 - A collection of objects, each described by its attributes,
and labeled with its classes

Supervised ML / Classification

- Apple vs Orange



Supervised ML / Classification


- Apple vs Orange



\neq

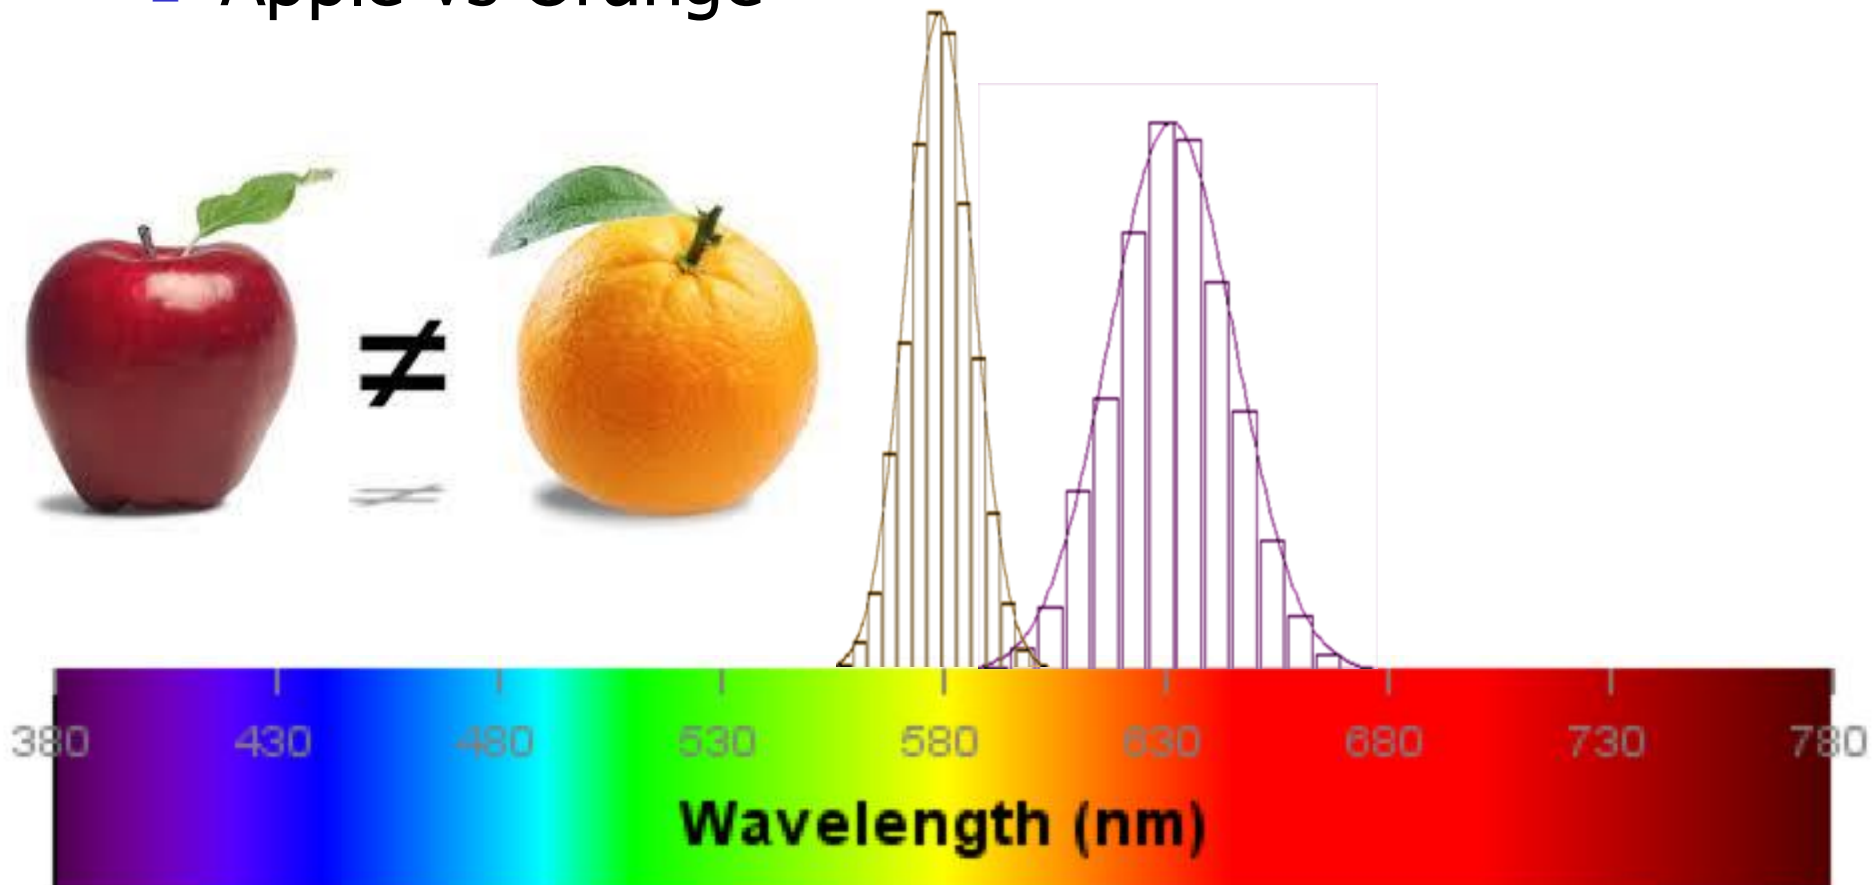


- Input Feature: object wavelength
- If wavelength > 600nm then apple
otherwise orange



Supervised ML / Classification


- Apple vs Orange



- Probability distribution of wavelength for Apples
- Probability distribution of wavelength for Oranges

Supervised ML / Classification

- Apple vs Orange
 - Probability distribution of wavelength for Apples
 - Probability distribution of wavelength for Oranges
- or:
 - Joint distribution over 2D space
 $(x,y) = (\text{wavelength}, \text{fruit class})$




Supervised ML / Classification

- Apple vs Orange

- Probability distribution of wavelength for Apples
- Probability distribution of wavelength for Oranges

or:

- Joint distribution over 2D space
 $(x,y) = (\text{wavelength}, \text{fruit class})$



- We often think in terms of either joint probability distribution $p(x,y)$, or probability distribution within each individual class, i.e., conditional probability of x given class, $p(x|y_i)$
- But what we really want for classification is $p(y_i|x)$:
what is the probability of class y_i (of apples, or of oranges) for given value of x ?


Supervised ML / Classification

- Apple vs Orange

- Probability distribution of wavelength for Apples
- Probability distribution of wavelength for Oranges

or:

- Joint distribution over 2D space
 $(x,y) = (\text{wavelength}, \text{fruit class})$




- For today, let just focus on predicting whether $p(\text{apple}|x) > 50\%$ (i.e., should we predict apple, or orange)

- That is, we will try to predict class y given wavelength x ,

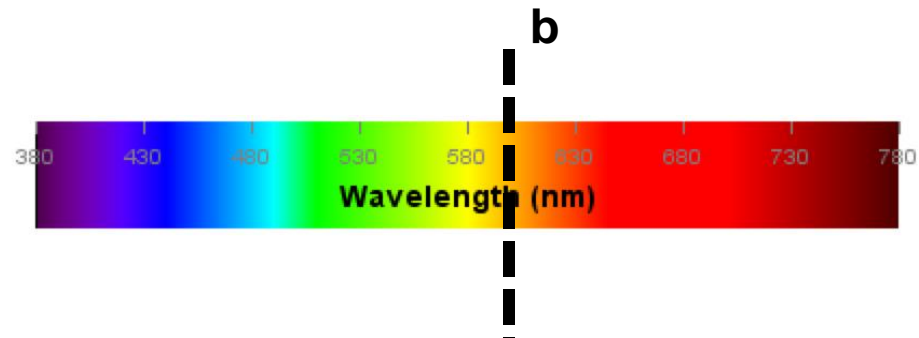
not the specific value of probability of that class

Supervised ML / Classification




- Apple vs Orange
- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
 - b is **unknown**, need to be learned from examples



What is the shape/plot
of function
 $f(x)=\text{sign}(x - b)$?




Supervised ML / Classification



- Apple vs Orange
- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x) = \text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
 - b is **unknown**, need to be learned from examples




Supervised ML / Classification




■ Apple vs Orange

- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x) = \text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
 - b is **unknown**, need to be learned from examples




An apple: $y=1$
predicted as
apple: $f(x)=1$

Supervised ML / Classification




■ Apple vs Orange

- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x) = \text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
 - b is **unknown**, need to be learned from examples




An orange: $y=1$
predicted as
apple: $f(x)=1$

Supervised ML / Classification




■ Apple vs Orange

- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x) = \text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
 - b is **unknown**, need to be learned from examples




An orange: $y = -1$
predicted as
orange: $f(x) = -1$

Supervised ML / Classification




■ Apple vs Orange

- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x) = \text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
 - b is **unknown**, need to be learned from examples




An apple: $y=1$
predicted as
orange: $f(x)=-1$

Supervised ML / Classification



■ Apple vs Orange

- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x) = \text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)



← this "b" ?


← or this "b" ?

← or maybe this "b" ?

- **b is unknown**, need to be learned from **examples**


different **b** => different **f(x)** => different predictions for the same **x**

Supervised ML / Classification




- Apple vs Orange
- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
 - The function has one input (x) and one trainable parameter (b)
- Algorithm:
 - Set initial value of trainable parameter b (0, or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)

Supervised ML / Classification




- Apple vs Orange
- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
- Algorithm:
 - Set initial value of b (0, or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update trainable parameter b
 - **How?**


Supervised ML / Classification




- Apple vs Orange
- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x) = \text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
- Algorithm:
 - Set initial value of b (0, or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i) = \text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is wrong, update b
 - Assume $y_i = 1$ (apple), but prediction is $f(x_i) = -1$ (orange)
 - Which of the four situations on the plot is it?




Supervised ML / Classification




- Apple vs Orange
- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x) = \text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
- Algorithm:
 - Set initial value of b (0, or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i) = \text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is wrong, update b
 - if $y_i = 1$ (apple),
but prediction is $f(x_i) = -1$ (orange),
how should we tweak " b "?




Supervised ML / Classification




- Apple vs Orange
- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x) = \text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
- Algorithm:
 - Set initial value of b (0, or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i) = \text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is wrong, update b
 - if $y_i = 1$ (apple), but prediction is $f(x_i) = -1$ (orange), how should we tweak " b "?
 - We need to move the threshold to the left. That is, **decrease b** . This may **increase $f()$** for that sample.




Supervised ML / Classification




- Apple vs Orange
- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x) = \text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
- Algorithm:
 - Set initial value of b (0, or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i) = \text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is wrong, update b
 - Assume $y_i = -1$ (orange), but prediction is $f(x_i) = +1$ (apple)
 - Which of the four situations on the plot is it?




Supervised ML / Classification



- Apple vs Orange
- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x) = \text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
- Algorithm:
 - Set initial value of b (0, or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i) = \text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is wrong, update b
 - Assume $y_i = -1$ (orange), but prediction is $f(x_i) = +1$ (apple) how should we tweak " b "?




Supervised ML / Classification




■ Apple vs Orange

- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x) = \text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
- Algorithm:
 - Set initial value of b (0, or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i) = \text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is wrong, update b
 - Assume $y_i = -1$ (orange), but prediction is $f(x_i) = +1$ (apple) how should we tweak " b "?
 - We need to move the threshold to the right. That is, increase b . This may decrease $f()$ for that sample.




Supervised ML / Classification



- Apple vs Orange
- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0: tough to predict)
- Algorithm:
 - Set initial value of b (0, or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1, f(x_i)=-1$, **decrease b** to **increase f**
 - if $y_i = -1, f(x_i)=1$, **increase b** to **decrease f**



Supervised ML / Classification

- Example run
- Start by setting $b=560$
- Each increase/decrease is by 10
-  $y=-1, x=580, f(x)=+1$, increase b to $b=570$



- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0 : tough to predict)
- Algorithm:
 - Set initial value of b (0 , or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1, f(x_i)=-1$, **decrease b to increase f**
 - if $y_i = -1, f(x_i)=1$, **increase b to decrease f**




Supervised ML / Classification

- Example run
 - Start by setting $b=560$
 - Each increase/decrease is by 10
- 
- $y=-1$, $x=580$, $f(x)=+1$, increase b to $b=570$
- 
- $y=+1$, $x=640$, $f(x)=+1$, do nothing ($b=570$)



- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0 : tough to predict)
- Algorithm:
 - Set initial value of b (0 , or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1$, $f(x_i)=-1$, **decrease b to increase f**
 - if $y_i = -1$, $f(x_i)=1$, **increase b to decrease f**

Supervised ML / Classification

- Example run
 - Start by setting $b=560$
 - Each increase/decrease is by 10
- 
 - $y=-1, x=580, f(x)=+1$, increase b to $b=570$
- 
 - $y=+1, x=640, f(x)=+1$, do nothing ($b=570$)
- 
 - $y=-1, x=580, f(x)=+1$, increase b to $b=580$



- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0 : tough to predict)
- Algorithm:
 - Set initial value of b (0 , or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1, f(x_i)=-1$, **decrease b to increase f**
 - if $y_i = -1, f(x_i)=1$, **increase b to decrease f**

Supervised ML / Classification

- Example run
- Start by setting $b=560$
- Each increase/decrease is by 10



- $y=-1, x=580, f(x)=+1$, increase b to $b=570$



- $y=+1, x=640, f(x)=+1$, do nothing ($b=570$)



- $y=-1, x=580, f(x)=+1$, increase b to $b=580$



- $y=-1, x=615, f(x)=+1$, increase b to $b=590$



- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0 : tough to predict)
- Algorithm:
 - Set initial value of b (0 , or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1, f(x_i)=-1$, **decrease b to increase f**
 - if $y_i = -1, f(x_i)=1$, **increase b to decrease f**

Supervised ML / Classification

- Example run
- Start by setting $b=560$
- Each increase/decrease is by 10



- $y=-1, x=580, f(x)=+1$, increase b to $b=570$



- $y=+1, x=640, f(x)=+1$, do nothing ($b=570$)



- $y=-1, x=580, f(x)=+1$, increase b to $b=580$



- $y=-1, x=615, f(x)=+1$, increase b to $b=590$









- $y=-1, x=615, f(x)=+1$, increase b to $b=600$



- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0 : tough to predict)
- Algorithm:
 - Set initial value of b (0 , or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1, f(x_i)=-1$, **decrease b to increase f**
 - if $y_i = -1, f(x_i)=1$, **increase b to decrease f**

Supervised ML / Classification

- Example run
- Start by setting $b=560$
- Each increase/decrease is by 10








-  $y=-1, x=580, f(x)=+1$, increase b to $b=570$
-  $y=+1, x=640, f(x)=+1$, do nothing ($b=570$)
-  $y=-1, x=580, f(x)=+1$, increase b to $b=580$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=590$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=600$
-  $y=+1, x=620, f(x)=+1$, do nothing ($b=600$)



- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0 : tough to predict)
- Algorithm:
 - Set initial value of b (0 , or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1, f(x_i)=-1$, **decrease b to increase f**
 - if $y_i = -1, f(x_i)=1$, **increase b to decrease f**

Supervised ML / Classification

- Example run
- Start by setting $b=560$
- Each increase/decrease is by 10

-  $y=-1, x=580, f(x)=+1$, increase b to $b=570$
-  $y=+1, x=640, f(x)=+1$, do nothing ($b=570$)
-  $y=-1, x=580, f(x)=+1$, increase b to $b=580$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=590$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=600$
-  $y=+1, x=620, f(x)=+1$, do nothing ($b=600$)
-  $y=-1, x=580, f(x)=-1$, do nothing ($b=600$)



- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0 : tough to predict)
- Algorithm:
 - Set initial value of b (0 , or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1, f(x_i)=-1$, **decrease b to increase f**
 - if $y_i = -1, f(x_i)=1$, **increase b to decrease f**

Supervised ML / Classification

- Example run

- Start by setting $b=560$

- Each increase/decrease is by 10



- $y=-1, x=580, f(x)=+1$, increase b to $b=570$



- $y=+1, x=640, f(x)=+1$, do nothing ($b=570$)



- $y=-1, x=580, f(x)=+1$, increase b to $b=580$



- $y=-1, x=615, f(x)=+1$, increase b to $b=590$



- $y=-1, x=615, f(x)=+1$, increase b to $b=600$



- $y=+1, x=620, f(x)=+1$, do nothing ($b=600$)



- $y=-1, x=580, f(x)=-1$, do nothing ($b=600$)



- $y=-1, x=615, f(x)=+1$, increase b to $b=610$



- Let:

- x denote object's **color** (wavelength in nm)
- y denote object's **class** (-1 = orange, 1 = apple)

- Prediction will be made by evaluating a function:

- $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0 : tough to predict)

- Algorithm:

- Set initial value of b (0 , or random, or a guess)










- Loop:

- Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
- Compare true class y_i with predicted class $f(x_i)$
- If prediction is right, go to next sample ($i=i+1$)
- If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1, f(x_i)=-1$, **decrease b to increase f**
 - if $y_i = -1, f(x_i)=1$, **increase b to decrease f**

Supervised ML / Classification

- Example run
- Start by setting $b=560$
- Each increase/decrease is by 10













-  $y=-1, x=580, f(x)=+1$, increase b to $b=570$
-  $y=+1, x=640, f(x)=+1$, do nothing ($b=570$)
-  $y=-1, x=580, f(x)=+1$, increase b to $b=580$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=590$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=600$
-  $y=+1, x=620, f(x)=+1$, do nothing ($b=600$)
-  $y=-1, x=580, f(x)=-1$, do nothing ($b=600$)
-  $y=-1, x=615, f(x)=+1$, increase b to $b=610$
-  $y=+1, x=640, f(x)=+1$, do nothing ($b=610$)

- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0 : tough to predict)
- Algorithm:
 - Set initial value of b (0 , or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1, f(x_i)=-1$, **decrease b to increase f**
 - if $y_i = -1, f(x_i)=1$, **increase b to decrease f**

Supervised ML / Classification

- Example run
- Start by setting $b=560$
- Each increase/decrease is by 10



-  $y=-1, x=580, f(x)=+1$, increase b to $b=570$
-  $y=+1, x=640, f(x)=+1$, do nothing ($b=570$)
-  $y=-1, x=580, f(x)=+1$, increase b to $b=580$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=590$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=600$
-  $y=+1, x=620, f(x)=+1$, do nothing ($b=600$)
-  $y=-1, x=580, f(x)=-1$, do nothing ($b=600$)
-  $y=-1, x=615, f(x)=+1$, increase b to $b=610$
-  $y=+1, x=640, f(x)=+1$, do nothing ($b=610$)
-  $y=+1, x=618, f(x)=+1$, do nothing ($b=610$)

- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0 : tough to predict)
- Algorithm:
 - Set initial value of b (0 , or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1, f(x_i)=-1$, **decrease b to increase f**
 - if $y_i = -1, f(x_i)=1$, **increase b to decrease f**












Supervised ML / Classification



- Example run

- Start by setting $b=560$

- Each increase/decrease is by 10

-  $y=-1, x=580, f(x)=+1$, increase b to $b=570$
-  $y=+1, x=640, f(x)=+1$, do nothing ($b=570$)
-  $y=-1, x=580, f(x)=+1$, increase b to $b=580$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=590$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=600$
-  $y=+1, x=620, f(x)=+1$, do nothing ($b=600$)
-  $y=-1, x=580, f(x)=-1$, do nothing ($b=600$)
-  $y=-1, x=615, f(x)=+1$, increase b to $b=610$
-  $y=+1, x=640, f(x)=+1$, do nothing ($b=610$)
-  $y=+1, x=618, f(x)=+1$, do nothing ($b=610$)
-  $y=-1, x=615, f(x)=+1$, increase b to $b=620$













- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0 : tough to predict)
- Algorithm:
 - Set initial value of b (0 , or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1, f(x_i)=-1$, **decrease b to increase f**
 - if $y_i = -1, f(x_i)=1$, **increase b to decrease f**

Supervised ML / Classification

- Example run

- Start by setting $b=560$

- Each increase/decrease is by 10

-  $y=-1, x=580, f(x)=+1$, increase b to $b=570$
-  $y=+1, x=640, f(x)=+1$, do nothing ($b=570$)
-  $y=-1, x=580, f(x)=+1$, increase b to $b=580$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=590$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=600$
-  $y=+1, x=620, f(x)=+1$, do nothing ($b=600$)
-  $y=-1, x=580, f(x)=-1$, do nothing ($b=600$)
-  $y=-1, x=615, f(x)=+1$, increase b to $b=610$
-  $y=+1, x=640, f(x)=+1$, do nothing ($b=610$)
-  $y=+1, x=618, f(x)=+1$, do nothing ($b=610$)
-  $y=-1, x=615, f(x)=+1$, increase b to $b=620$
-  $y=+1, x=618, f(x)=-1$, decrease b to $b=610$



- Let:

- x denote object's **color** (wavelength in nm)
- y denote object's **class** (-1 = orange, 1 = apple)

- Prediction will be made by evaluating a function:

- $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0 : tough to predict)

- Algorithm:

- Set initial value of b (0 , or random, or a guess)
- Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1, f(x_i)=-1$, **decrease b to increase f**
 - if $y_i = -1, f(x_i)=1$, **increase b to decrease f**














Supervised ML / Classification



- Example run

- Start by setting $b=560$

- Each increase/decrease is by 10

-  $y=-1, x=580, f(x)=+1$, increase b to $b=570$
-  $y=+1, x=640, f(x)=+1$, do nothing ($b=570$)
-  $y=-1, x=580, f(x)=+1$, increase b to $b=580$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=590$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=600$
-  $y=+1, x=620, f(x)=+1$, do nothing ($b=600$)
-  $y=-1, x=580, f(x)=-1$, do nothing ($b=600$)
-  $y=-1, x=615, f(x)=+1$, increase b to $b=610$
-  $y=+1, x=640, f(x)=+1$, do nothing ($b=610$)
-  $y=+1, x=618, f(x)=+1$, do nothing ($b=610$)
-  $y=-1, x=615, f(x)=+1$, increase b to $b=620$
-  $y=+1, x=618, f(x)=-1$, decrease b to $b=610$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=620$

- Let:

- x denote object's **color** (wavelength in nm)
- y denote object's **class** (-1 = orange, 1 = apple)

- Prediction will be made by evaluating a function:

- $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0 : tough to predict)

- Algorithm:

- Set initial value of b (0 , or random, or a guess)

- Loop:

- Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
- Compare true class y_i with predicted class $f(x_i)$
- If prediction is right, go to next sample ($i=i+1$)
- If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1, f(x_i)=-1$, **decrease b to increase f**
 - if $y_i = -1, f(x_i)=1$, **increase b to decrease f**















Supervised ML / Classification



- Example run

- Start by setting $b=560$

- Each increase/decrease is by 10

-  $y=-1, x=580, f(x)=+1$, increase b to $b=570$
-  $y=+1, x=640, f(x)=+1$, do nothing ($b=570$)
-  $y=-1, x=580, f(x)=+1$, increase b to $b=580$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=590$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=600$
-  $y=+1, x=620, f(x)=+1$, do nothing ($b=600$)
-  $y=-1, x=580, f(x)=-1$, do nothing ($b=600$)
-  $y=-1, x=615, f(x)=+1$, increase b to $b=610$
-  $y=+1, x=640, f(x)=+1$, do nothing ($b=610$)
-  $y=+1, x=618, f(x)=+1$, do nothing ($b=610$)
-  $y=-1, x=615, f(x)=+1$, increase b to $b=620$
-  $y=+1, x=618, f(x)=-1$, decrease b to $b=610$
-  $y=-1, x=615, f(x)=+1$, increase b to $b=620$
-  $y=+1, x=618, f(x)=-1$, decrease b to $b=610$
- ...

- Let:

- x denote object's **color** (wavelength in nm)
- y denote object's **class** (-1 = orange, 1 = apple)


- Prediction will be made by evaluating a function:

- $f(x)=\text{sign}(x - b)$ - it returns either -1 or 1 (or 0 : tough to predict)

- Algorithm:


- Set initial value of b (0 , or random, or a guess)
- Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - In a way that would make it more likely to get correct prediction for the current sample
 - if $y_i = 1, f(x_i)=-1$, **decrease b to increase f**
 - if $y_i = -1, f(x_i)=1$, **increase b to decrease f**

Supervised ML / Classification




- Apple vs Orange
- Algorithm:
 - Set initial value of b (0, or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i) = \text{sign}(x_i - b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - if $y_i = 1, f(x_i) = -1$, what should we do? **Decrease b** to increase f
 - $y_i - f(x_i) = 2$
 - $b_{\text{new}} = b_{\text{old}} - 2c$
 - if $y_i = -1, f(x_i) = 1$, what should we do? **Increase b** to decrease f
 - $y_i - f(x_i) = -2$
 - $b_{\text{new}} = b_{\text{old}} + 2c$
 - **What is “ c ”? The “learning rate”, a small number chosen by the user**
 - Single formula: $b_{\text{new}} = b_{\text{old}} - c[y_i - f(x_i)]$

Supervised ML / Classification




- Apple vs Orange
- Algorithm (with small tweak, sign in front of b):
 - Set initial value of b (0, or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i) = \text{sign}(x_i + b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - if $y_i = 1, f(x_i) = -1$, what should we do? **Increase b** to increase f
 - $y_i - f(x_i) = 2$
 - $b_{\text{new}} = b_{\text{old}} + 2c$
 - if $y_i = -1, f(x_i) = 1$, what should we do? **Decrease b** to decrease f
 - $y_i - f(x_i) = -2$
 - $b_{\text{new}} = b_{\text{old}} - 2c$
 - Single formula: $b_{\text{new}} = b_{\text{old}} + c[y_i - f(x_i)]$

Supervised ML / Classification




- Apple vs Orange
- Algorithm
 - unified, w_0 can be set to -1 (original version) or +1 (tweaked version):
 - Set initial value of b (0, or random, or a guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i) = \text{sign}(x_i + w_0 b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - $b_{\text{new}} = b_{\text{old}} + c[y_i - f(x_i)]w_0$

Supervised ML / Classification



- Apple vs Orange
- We have a simple predictive model that involves one trainable parameter, the threshold value separating apples from oranges
- We have a training algorithm for that model
- **What other ways are there to specify trainable parameters?**


Supervised ML / Classification



- Apple vs Orange
- We have a simple predictive model that involves one trainable parameter, the threshold value separating apples from oranges
- We have a training algorithm for that model
- **What other ways are there to specify trainable parameters?**
 - We can add a parameter that scales the input x

Supervised ML / Classification

- Using a parameter to scale an input attribute
- Let:
 - x denote object's **color** (wavelength in nm)
 - y denote object's **class** (-1 = orange, 1 = apple)
- Prediction will be made by evaluating a function:
 - $f(x)=\text{sign}(wx - 1)$ - it returns either -1 or 1 (or 0: tough to predict)
 - w is **unknown**, needs to be learned from examples



- E.g. if wavelength 590 separates apples/oranges, then $w=1/590$ is a good choice
 - $x=600$ will lead to $wx-1 = 600/590 - 1 = 1.0169 - 1$ $f>0(\text{apple})$
 - $x=580$ will lead to $wx-1 = 580/590 - 1 = 0.9831 - 1$ $f<0(\text{orange})$



Supervised ML / Classification

- Using a parameter to scale an input attribute
- Algorithm
 - Set initial values of w (random guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i) = \text{sign}(wx_i - 1)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update w
 - $w_{\text{new}} = w_{\text{old}} + c[y_i - f(x_i)]x_i$
 - Why x_i ?
 - $y=+1, f(x)<0, [y-f(x)]>0$, we need to increase $f(x)$
 - $x>0$: to increase $f(x)$, increase w
 - $x<0$: to increase $f(x)$, decrease w
 - $y=-1, f(x)>0, [y-f(x)]<0$, we need to decrease $f(x)$
 - $x>0$: to decrease $f(x)$, decrease w
 - $x<0$: to decrease $f(x)$, increase w

We could just use: $w_{\text{new}} = w_{\text{old}} + c[y_i - f(x_i)]\text{sign}(x_i)$
but using x_i also helps us tweak magnitude of the change



Supervised ML / Classification

- Using a parameter to scale an input attribute
- Prediction will be made by evaluating a function:
 - $f(x) = \text{sign}(wx - 1)$ - it returns either -1 or 1 (or 0: tough to predict)
 - should return the same prediction as
 - $f(x) = \text{sign}(x - 1/w)$



**With one attribute (here, wavelength),
training the scale w , or training the threshold b
is essentially the same***

Supervised ML / Classification

- Using a parameter to scale an input attribute
 - Prediction will be made by evaluating a function:
 - $f(x) = \text{sign}(wx - 1)$ - it returns either -1 or 1 (or 0: tough to predict)
 - should return the same prediction as
 - $f(x) = \text{sign}(x - 1/w)$


**With one attribute (here, wavelength),
training the scale w , or training the threshold b
is essentially the same***

* one small difference: with threshold b , there's no way to achieve inverted predictions, i.e., if we set apples -1, oranges +1, no threshold will work




Supervised ML / Classification

- Apple vs Orange beyond single threshold



\neq





- Feature A: object wavelength
- If $A > 600$ apple
otherwise: if $A < 560$ apple
otherwise orange



Supervised ML / Classification

- Apple vs Orange beyond single threshold



- Feature A: object wavelength
- If $A > 600$ apple
otherwise: if $A < 560$ apple
otherwise ???






WHAT SHOULD WE DO?

Wavelength alone is not enough






Any other attributes/features?

Multiple attributes / features



- **Feature A: color**
 - Apple: green, red, orange
 - Orange: orange
- **Feature B: color variability**
 - Apple: uniform or not
 - Orange: uniform
- **Feature C: texture**
 - Apple: smooth
 - Orange: rough
- **Feature D: reflectance**
 - Apple: reflects more light
 - Orange: reflects less light
- **Feature E: shape**
 - Apple: non-convex
 - Orange: convex (almost)

Multiple attributes / features



- **Feature A: color**
 - Apple: green, red, orange
 - Orange: orange
- **Feature B: color variability**
 - Apple: uniform or not
 - Orange: uniform
- **Feature C: texture**
 - Apple: smooth
 - Orange: rough
- **Feature D: reflectance**
 - Apple: reflects more light
 - Orange: reflects less light
- **Feature E: shape**
 - Apple: non-convex
 - Orange: convex (almost)
- Using trainable parameters (“w’s”) to scale individual features becomes very useful when we have >1 feature:
 - we can adjust their importance
 - we can adjust their sign (i.e., high means apple or high means orange)

Supervised ML: Perceptron

- One way of making predictions with samples that have 2 features, $x=(x^1, x^2)$

- Define 2 trainable parameters: weights w_1 and w_2

- Prediction is: $f(x_i)=\text{sign}(w_1x_i^1 + w_2x_i^2)$

Note: here, ² is not 2nd power, just another index:
 $x_i^2=x[i][2]$

- Perceptron Algorithm

- Set initial values of w_1 and w_2 (random guess)

- Loop:

- Present a sample x_i and predict $f(x_i)=\text{sign}(w_1x_i^1 + w_2x_i^2)$

- Compare true class y_i with predicted class $f(x_i)$

- If prediction is right, go to next sample ($i=i+1$)

- If prediction is wrong, update w :

- $w_{1,\text{new}}=w_{1,\text{old}} + c[y_i - f(x_i)]x_i^1$

- $w_{2,\text{new}}=w_{2,\text{old}} + c[y_i - f(x_i)]x_i^2$

Same as we did before with one “w”,
but now it’s done separately to each “w_j”
using corresponding feature x_i^j

Same, using vectors

- One way of making predictions with samples that have 2 features, each sample i is a 2D vector: $\mathbf{x}_i = (x_i^1, x_i^2)$
- Define a 2D vector $\mathbf{w} = (w_1, w_2)$ of 2 trainable parameters

- Prediction is: $f(\mathbf{x}_i) = \text{sign}(\mathbf{w}_1 x_i^1 + \mathbf{w}_2 x_i^2) = \text{sign}(\mathbf{w}^T \mathbf{x}_i)$
- Vector (2x1) \rightarrow \mathbf{w}^T is transpose op., makes it (1x2)
- Matrix multiplication: (1x2) times (2x1) \rightarrow Vector (2x1)

$$\mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} \leftarrow \text{2D Vector (2x1, a column vector)}$$

$$\mathbf{x} = \begin{bmatrix} x^1 \\ x^2 \end{bmatrix} \leftarrow \text{2D Vector (2x1, a column vector)}$$

$$\mathbf{w}^T = [w_1, w_2] \leftarrow \text{Transposed vector (1x2)}$$

Matrix multiplication:
(1x2) times (2x1)
gives a
single number (1x1)

$$\mathbf{w}^T \mathbf{x} = [w_1, w_2] \begin{bmatrix} x^1 \\ x^2 \end{bmatrix} = w_1 x^1 + w_2 x^2$$

Perceptron using vectors

- One way of making predictions with samples that have 2 features, $x=(x^1, x^2)$
 - Define a 2D vector $w=(w_1, w_2)$ of 2 trainable parameters

- Prediction is: $f(x_i)=\text{sign}(w_1x_i^1 + w_2x_i^2) = \text{sign}(w^T x)$
 - Vector (2x1) → w
 - T is transpose op., makes it (1x2)
 - Vector (2x1) → x
 - Matrix multiplication: (1x2) times (2x1)

- Perceptron Algorithm
 - Set initial values of vector w (random guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(w^T x)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong (incl. $f(x_i)=0$, no prediction), update w :
 - $w_{\text{new}} = w_{\text{old}} + c[y_i - f(x_i)]x_i$
 - vector → w_{new}
 - vector → w_{old}
 - vector → x_i



HW1 preview

- HW1 will be announced after add/drop, on Wednesday (8/28)
- You will have until 9/10, 5pm to complete it
- It will involve a simple dataset with 2 classes and 2 features
- Your task will be to explore the process of training a perceptron (the vector version) from previous lecture slide
- Python libraries to be used:
 - Pandas (reading in a csv file)
 - Matplotlib (plotting diagrams of training progress)
 - Numpy (storing vectors, doing the math with them)
 - **ML libraries (e.g. sklearn, pytorch, tensorflow, others) not allowed**
- The underlying goal of this simple HW is to bring everyone up to speed with python and its basic libraries for routine ML tasks like reading input, plotting, etc.
- If you are not confident with python, practice it before next Wednesday





Summary

- What we have seen is supervised learning:
 - There is a true class (e.g. type of fruit) that is unknown and should be predicted
- Supervised learning is a major type of machine learning
 - Includes self-supervised learning where the true class comes from the input features
- Other important types of ML are:
 - Unsupervised learning (e.g. cluster objects together)
 - Reinforcement learning (learn to choose good action, with distant reward, e.g. chess)