


Introduction to Machine Learning

Lecture 3

Instructor:
Dr. Tom Arodz

Recap: Simple threshold-based predictive model for single feature



- **Data:** Two classes (y_i is -1 or +1), one feature (x_i is a real number)
- Predictive model:** $f(x_i) = \text{sign}(x_i + w_0 b)$
with one trainable parameter b
 w_0 can be set to -1 (some textbooks) or +1 (other textbooks)


- Set initial value of b (0, or random, or a guess)
- Loop:
 - Present a sample x_i and predict $f(x_i) = \text{sign}(x_i + w_0 b)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update b
 - $b_{\text{new}} = b_{\text{old}} + c[y_i - f(x_i)]w_0$

Learning rate (small positive number chosen by user, e.g., 0.001)

e.g., if we want $f(x_i)$ to increase from -1 to +1 because y_i is +1:

- increase b if $w_0 = +1$
- decrease b if $w_0 = -1$

Recap: Supervised ML





- Typically one feature is not sufficient




- We have many features

- Possible features

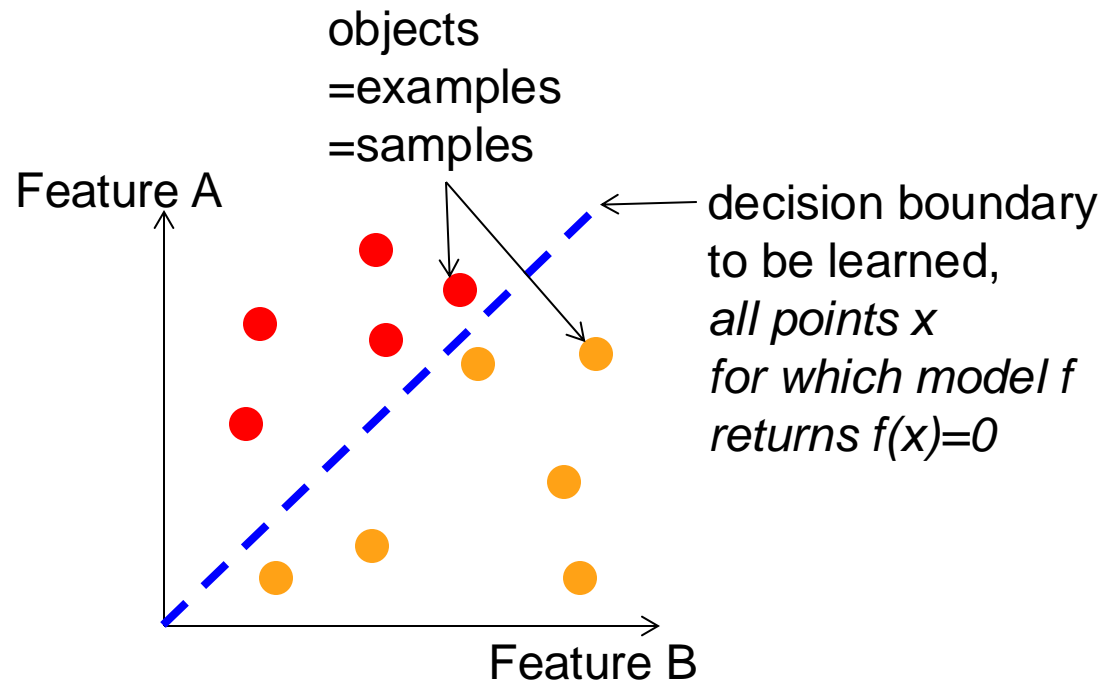
- color
- color variability
- texture
- reflectance
- shape



Recap: Supervised ML



- Building a **classifier (a predictive model)**
 - F-dimensional feature space: R^F
 - Two features $\Rightarrow F=2$
 - Training a classifier can be seen as:
 - finding a **decision boundary** between **class 1** and **class -1**



Recap: Simple predictive model for two features

- Data: Two classes (y_i is -1 or +1),
two features (real numbers x_i^1 and x_i^2)

Predictive model: $f(\mathbf{x}_i) = \text{sign}(\mathbf{w}^T \mathbf{x}) = \text{sign}(w_1 x_i^1 + w_2 x_i^2)$
with two trainable parameters w_1 and w_2

- Set initial value of w_1 and w_2 (random, or a guess)
- Loop:
 - Present a sample \mathbf{x}_i and predict $f(\mathbf{x}_i) = \text{sign}(w_1 x_i^1 + w_2 x_i^2)$
 - Compare true class y_i with predicted class $f(\mathbf{x}_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update \mathbf{w}
 - $\mathbf{w}_{\text{new}} = \mathbf{w}_{\text{old}} + c[y_i - f(\mathbf{x}_i)]\mathbf{x}_i$
 - that is:
 - $w_{1,\text{new}} = w_{1,\text{old}} + c[y_i - f(\mathbf{x}_i)]x_i^1$
 - $w_{2,\text{new}} = w_{2,\text{old}} + c[y_i - f(\mathbf{x}_i)]x_i^2$

e.g., $y_i=+1$, $\mathbf{x}_i=(0.75, -0.25)$, $w_1=w_2=-1$: we want $f(\mathbf{x}_i)$ to increase from -1 to +1:
increase w_1 (from negative -1 towards >0 , so that $w_1 x^1$ becomes positive)
but decrease w_2 (make it more negative, so that $w_2 x^2$ is larger positive)

Recap: Simple predictive model for two features

- Data: Two classes (y_i is -1 or +1),
two features (real numbers x_i^1 and x_i^2)

Predictive model: $f(\mathbf{x}_i) = \text{sign}(\mathbf{w}^T \mathbf{x}) = \text{sign}(w_1 x_i^1 + w_2 x_i^2)$
with two trainable parameters w_1 and w_2

- Set initial value of w_1 and w_2 (random, or a guess)
- Loop:
 - Present a sample \mathbf{x}_i and predict $f(\mathbf{x}_i) = \text{sign}(w_1 x_i^1 + w_2 x_i^2)$
 - Compare true class y_i with predicted class $f(\mathbf{x}_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update \mathbf{w}
 - $\mathbf{w}_{\text{new}} = \mathbf{w}_{\text{old}} + c[y_i - f(\mathbf{x}_i)]\mathbf{x}_i$
 - that is:
 - $w_{1,\text{new}} = w_{1,\text{old}} + c[y_i - f(\mathbf{x}_i)]x_i^1$
 - $w_{2,\text{new}} = w_{2,\text{old}} + c[y_i - f(\mathbf{x}_i)]x_i^2$

e.g., $y_i = -1$, $\mathbf{x}_i = (0.75, 0.25)$, $w_1 = w_2 = 1$: we want $f(\mathbf{x}_i)$ to decrease from +1 to -1:
decrease w_1 and w_2 ,
but decrease w_1 faster (by $c \cdot 2 \cdot 0.75$)
and w_2 slower (by $c \cdot 2 \cdot 0.25$)

Why? w_1 gets multiplied by larger number ($x^1=0.75$) so we want to decrease it more than w_2

Recap: Perceptron using vectors

- One way of making predictions with samples that have 2 features, $x=(x^1, x^2)$
 - Define a 2D vector $w=(w_1, w_2)$ of 2 trainable parameters

- Prediction is: $f(x_i)=\text{sign}(w_1x_i^1 + w_2x_i^2) = \text{sign}(w^T x)$
 - Vector (2x1) → w
 - T is transpose op., makes it (1x2)
 - Vector (2x1) → x
 - Matrix multiplication: (1x2) times (2x1)

- Perceptron Algorithm
 - Set initial values of vector w (random guess)
 - Loop:
 - Present a sample x_i and predict $f(x_i)=\text{sign}(w^T x)$
 - Compare true class y_i with predicted class $f(x_i)$
 - If prediction is right, go to next sample ($i=i+1$)
 - If prediction is wrong, update w

- $w_{\text{new}} = w_{\text{old}} + c[y_i - f(x_i)]x_i$


vector vector

vector



Supervised ML / Classification




- Building a ***Classifier***



- Feature generation: measure attributes, transform them if needed
- Feature selection (optional): select most relevant features
- Learning: train a machine learning model
- Validation: test the model, redo the above if not satisfied

Supervised ML / Classification


- Building a ***Classifier***



- Features are typically numerical (real numbers, or binary)
 - If not, some numerical encoding is typically used
- Feature space:
 - if we have 2 features for each object
 - then each object is a point in a 2D space
 - 1000 features => 1000 dimensions

Supervised ML / Classification

- Building a ***Classifier***



- F-dimensional feature space: \mathbb{R}^F
- If we have just two classes (A/B, apples/oranges,...)
- Then, building a classifier translates to:
 - assigning regions of feature space to class A and the rest to class B
 - finding a ***decision boundary*** in feature space between class A and class B



Summary


- What kind of learning are we talking about?
 - Learning from experience/examples to predict to which of two classes an object belongs
- What are features?
 - The descriptors/attributes (typically numerical) that describe each object/sample
 - Each sample is a vector in the feature space
- What is a decision boundary?
 - A line/surface/hypersurface that has most of samples from one class on one side, and most samples from the other class on the other side
- What is a classifier? Why classification is “supervised machine learning”?
 - A model (mathematical formula) that predicts a class of an example based on its features
 - It’s supervised because the learning algorithm bases its actions on the availability of information about true classes of some samples
- How to build a classifier?
 - One simple way is to use a perceptron



- If you got a mixed breed dog and you don't know what the breeds are...
- ... you could consult an expert, who will look at:
 - size of the dog
 - fur
 - shape of tail, etc.,and compare these with what he/she learned about dog breeds
- It's learning to predict classes from features!
 - except, more than one class is correct

Supervised ML / Classification

- If you got a mixed breed dog and you don't know what the breeds are
- You could consult an expert, who will look at:
 - size of the dog
 - fur
 - shape of tail, etc.,and compare these with what he/she learned about dog breeds
- Or you could order a genetic test!
 - If you know the attribute that truly differentiates the classes
 - And you can measure it
 - Then there's nothing to learn!
 - Classification methods are useful if the differences between classes are not known or not easily quantified



A fully accurate feature eliminates the need for machine learning!



ML: Typical assumptions

- The features are somewhat informative but not perfectly correlated with the class
 - If features were perfect, no need for machine learning

Supervised ML / Classification

- Rain or sun tomorrow afternoon
 - Learning from experience:
 - *Ring around the moon?*
Rain real soon!
 - *Red sky at night, sailors delight.*
Red sky in morning, sailors take warning!



Supervised ML / Classification


- Rain or sun tomorrow afternoon

- Learning from experience:

- *Ring around the moon?
Rain real soon!*
- *Red sky at night, sailors delight.
Red sky in morning, sailors take warning!*



- Predictions from supercomputer simulation model
- There are other, better ways of predicting than machine learning
- Accurate simulation beats learning from data unless the problem is very complex





ML: Typical assumptions



- We do not have a detailed physical/mechanical understanding of the modeled object/process, or, if we do, the process is too complex to simulate
 - Otherwise we would apply known scientific formula or run a computer simulation to get our predictions

This further reinforces the assumption that:

- The features are somewhat informative but very strongly correlated with the class
 - If they were perfectly informative, no need for machine learning

Yet another ML example


Circle vs **other closed figure**



objects
=examples
=samples



Yet another ML example

- Circle vs other closed figure
 - Feature A: figure circumference
- Is this a useful/informative feature for this classification task?





Yet another ML example

- Circle vs other closed figure
 - Feature B: circumference of circumscribed circle
- Is this a useful/informative feature for this classification task?






Yet another ML example

- Circle vs other closed figure
 - Feature A: figure circumference
 - Feature B: circumference of circumscribed circle
- A,B uninformative individually, but informative together!
 - If $A < B$, not a circle
 - If $A = B$, a circle




Supervised ML / Classification

- Circle vs other closed figure
 - Feature A: figure circumference
 - Feature B: circumference of circumscribed circle
 - If $A < B$, not a circle (return -1)
 - If $A = B$, a circle (return 1)



Supervised ML / Classification



- Circle vs other closed figure
 - Feature A: figure circumference
 - Feature B: circumference of circumscribed circle
 - If $A < B$, not a circle
 - If $A = B$, a circle



?




Supervised ML / Classification

- Circle vs other closed figure
 - Feature A: figure circumference
 - Feature B: circumference of circumscribed circle
 - If $A < B$, not a circle
 - If $A = B$, a circle **FALSE if we allow non-convex shapes**



Supervised ML / Classification

- Circle vs other closed figure
 - Feature A: figure circumference
 - Feature B: circumference of circumscribed circle
 - If $A < B$, not a circle
 - If $A = B$, a circle **FALSE if non-convex**






ML: Typical assumptions

- We do not have a detailed physical/mechanical understanding of the modeled object/process, or the process is too complex
- The features are somewhat informative but not very strongly correlated with the class
- The association between features and class we can learn is likely to be accurate only for objects similar to our training set
 - We do not have much information about objects fundamentally different from we have seen during training

Supervised ML / Classification


- Gold vs fake
 - Magnetism
 - Not magnetic – but so is lead, etc.
 - Hardness
 - Soft – but so is lead, or lead covered with gold
 - Density = mass/vol.
 - Much denser than lead



Supervised ML / Classification

- Gold vs fake
 - Hardness
 - Soft – but so is anything covered with gold
 - Density = mass/vol.
 - Much denser than lead
 - But tungsten has the same density as gold
 - And is not magnetic

↑ Hardness



→ Density


Gold bar filled with cheap tungsten (wolfram) rods inside!

Hardness/density classifier will fail!

Supervised ML / Classification

- Gold vs fake
(also, is this cheap gold deal a spam?)

Hardness



Density



The classification problem changed over time


Classifier trained previously *may no longer work!*

Typically, we assume that the problem is stable:

- ***past and future sample are similar***

Supervised ML / Classification

- Gold vs fake
 - Heat capacity
 - Electric conductivity
- These can tell gold easily
(especially both features together)








ML: Typical assumptions

- We do not have a detailed physical/mechanical understanding of the modeled object/process, or the process is too complex
- The features are somewhat informative but not perfectly correlated with the class
- The association between features and class we can learn is likely to be accurate only for objects similar to our training set
- The association between *regions of feature space* and the *class variable* is fixed
 - If it was changing quickly over time, classifier would soon stop being accurate

ML: Typical assumptions

- What is the mathematical framework in which we can talk about imperfect, typically weak *associations* between *features/regions of feature space* and the *class variable*?

- **Probability!**



- Modern way of thinking about Machine Learning starts with framing learning in a probabilistic way
 - For next class, please revisit what you know about probability from previous courses



News & updates

- Homework 1 is out in Canvas
 - It is due on Thursday, 9/12 (in two weeks), at 5pm
- Next Monday (9/2) is Labor Day holiday, VCU will be closed, **no class**
- Next Wednesday (9/4) there will be **no class** (I have an appointment that I wasn't able to reschedule)
 - There WILL be office hours, starting at 1pm

News & updates

- Tentative schedule of topics, tests, homeworks is also in Canvas

Lectures			Homework assignments (+/- day or two)		
8/21/24	Wednesday	1 Course Intro (today)	8/21/24	Wednesday	
8/26/24	Monday	2 ML Example: Apples vs Oranges	8/26/24	Monday	
8/28/24	Wednesday	3 ML Assumptions / When to use ML	8/28/24	Wednesday	HW1 out
9/2/24	Monday	(Labor day)	9/2/24	Monday	
9/4/24	Wednesday	(cancelled)	9/4/24	Wednesday	
9/9/24	Monday	4 Probabilistic View of ML	9/9/24	Monday	
9/11/24	Wednesday	5 Maximum Likelihood Estimation	9/11/24	Wednesday	HW1 due (9/12)
9/16/24	Monday	6 Bayesian Learning	9/16/24	Monday	HW2 out
9/18/24	Wednesday	7 Trainable Probabilistic Models	9/18/24	Wednesday	
9/23/24	Monday	8 Error Metrics	9/23/24	Monday	
9/25/24	Wednesday	9 Loss Functions	9/25/24	Wednesday	
9/30/24	Monday	10 Cross-Entropy	9/30/24	Monday	HW2 due (10/1)
10/2/24	Wednesday	11 Intro to PyTorch, Test 1 Study Problems	10/2/24	Wednesday	
10/7/24	Monday	12 Test 1	10/7/24	Monday	
10/9/24	Wednesday	13 Nonlinear Kernel Methods	10/9/24	Wednesday	HW3 out
10/14/24	Monday	14 Nonlinear Features and Neural Network	10/14/24	Monday	
10/16/24	Wednesday	15 Details of an MLP Layer	10/16/24	Wednesday	
10/21/24	Monday	16 MLP: Universal Approximation	10/21/24	Monday	
10/23/24	Wednesday	17 MLP: Why Depth	10/23/24	Wednesday	
10/28/24	Monday	18 Training Deep MLPs	10/28/24	Monday	HW3 due
10/30/24	Wednesday	19 Convolutional Neural Networks	10/30/24	Wednesday	HW4 out
11/4/24	Monday	20 Weight Init, and Dropout	11/4/24	Monday	
11/6/24	Wednesday	21 Normalizations and Skip Connections	11/6/24	Wednesday	
11/11/24	Monday	22 LLM Intro	11/11/24	Monday	
11/13/24	Wednesday	23 Self-Attention and Transformers	11/13/24	Wednesday	HW5 out
11/18/24	Monday	24 Large Language Models	11/18/24	Monday	
11/20/24	Wednesday	25 Modern Models for Vision Task	11/20/24	Wednesday	
11/25/24	Monday	(Thanksgiving break)	11/25/24	Monday	
11/27/24	Wednesday	(Thanksgiving break)	11/27/24	Wednesday	
12/2/24	Monday	26 Diffusion-based Generative Models	12/2/24	Monday	
12/4/24	Wednesday	27 Summary, Test 2 Study Problems	12/4/24	Wednesday	HW5 due
12/9/24	Monday	28 Test 2	12/9/24	Monday	

- 
-
- Now, we move to HW1 slides
(they are in Canvas)